Pose-based Video Start Alignment on Mobile: Technical Note with Optional Audio-Assisted Refinement
Pose系列による動画の開始点自動同期（端末内）：音声支援の粗密化・検証を含む技術ノート

> **Open Technical Note (text-only).** This document is shared to help practitioners reproduce and improve the method, and to quietly establish prior art for the described techniques.

Abstract (EN)

This technical note describes an on-device method to estimate the temporal offset between two practice videos containing the same or highly similar performance segment (e.g., dance/ballet), enabling synchronized side-by-side playback.

The method samples frames at a fixed FPS and extracts pose keypoints per frame using an on-device pose estimator. It covers both single-person and multi-person pose estimators; for multi-person outputs, a target person can be selected using detection score, average keypoint confidence, spatial consistency, track continuity, and/or region-of-interest constraints.

Keypoints are normalized into a pose feature vector by filtering low-confidence points, selecting a robust center (pelvis-first, then shoulders, then centroid fallback), and applying scale normalization (shoulder width, hip width, or maximum radius fallback). Temporal alignment is obtained by scanning candidate offsets within ±searchRangeSec and computing an offset score as the average pose similarity across aligned frames. Per-frame similarity may be cosine similarity, and distance can be defined as 1−cosSim. Frame pairs with fewer than a minimum number of commonly valid keypoints (shared-valid points) are excluded from scoring to improve robustness.

Offset sign convention: Let Δ denote the shift applied to Video B relative to Video A when pairing frames for scoring. For a frame index *t* in A, the paired index in B is *(t + Δ)*. Thus, **Δ > 0** indicates that **B lags A** (B is evaluated later than A), and **Δ < 0** indicates that **B leads A**.

The selected offset may be converted into a robust start point using the earliest contiguous above-threshold similarity segment (or equivalent window-based criteria), with staged relaxation retries if no such segment is found.

A confidence score determines acceptance:

```text
confidence = 0.5·bestSim + 0.3·margin + 0.2·validRate
```

Here, margin is the gap between best and runner-up offset scores, and validRate denotes a validity ratio computed over the matched segment/window (e.g., valid-frame ratio and/or valid-keypoint ratio). When distance-minimization is used internally, bestSim and runner-up values for margin can be computed on any monotonic transform consistent with the acceptance logic (e.g., converting distance to similarity), without loss of generality.

Optional hybrids are also described:
- Audio-assisted coarse candidate generation (top-K offsets with minimum separation) followed by pose refinement within local neighborhoods, and/or
- Post-pose audio verification or micro-refinement around the selected offset to improve precision.

Additionally, an optional pose-only phase acceleration is described (e.g., restricting search to early segments or high-motion windows, with fallback to full search).

Keywords: pose estimation, skeleton keypoints, temporal alignment, time offset estimation, sequence matching, cosine similarity, contiguous match, confidence scoring, coarse-to-fine, audio-to-pose, audio verification, micro-refinement, on-device, mobile, multi-pose

---

概要 (JP)

本技術ノートは、同一または高類似の演技区間を含む練習動画2本（例：バレエ／ダンス）を左右に並べて比較する際に、両動画の開始位置（時間オフセット）を端末内で自動推定し、同期再生に用いる方式をまとめたものです。実務者が再現・改善できる情報提供を目的に公開します。併せて、記載手法を先行技術（prior art）として記録する意図も含みます。

手法は、一定FPSでフレーム列を作成し、各フレームで骨格キーポイント（座標＋信頼度）を推定する。単一人物推定に限らず、複数人物推定（Multi-Pose）の出力から対象人物を選別して用いる形態を含む（例：検出スコア、平均キーポイント信頼度、空間的一貫性、追跡の連続性、ROI制約など）。

低信頼キーポイントを除外した上で、中心（骨盤優先→肩→重心フォールバック）とスケール（肩幅／腰幅／最大半径フォールバック）により正規化した骨格ベクトル系列を構築する。次に、±searchRangeSec の範囲で時間オフセットを走査し、対応フレーム間の骨格ベクトル類似度（例：コサイン類似度）を平均集計したスコアで最良オフセットを選ぶ（距離は 1−cosSim と定義可能）。この際、両フレームで共通して有効なキーポイント数が所定数未満のフレーム対は、スコア計算対象から除外し頑健性を高める。

符号規約： Δ を「Video A に対して Video B 側を評価する時間（フレーム）シフト」とし、Aのフレームtに対しBのフレーム(t+Δ)を対応づける。Δ>0 は「BがAより遅れている（BはAより後の時刻として評価される）」、Δ<0 は「BがAより先行している」を意味する。

最良オフセットは、類似度が閾値以上となる連続区間（連続一致）を用いて頑健な開始点へ変換し、該当区間が見つからない場合は段階的に緩和して再試行する。

採否判定には以下の信頼度を用いる：

```text
confidence = 0.5×bestSim + 0.3×margin + 0.2×validRate
```

```



---

## 1. Problem Setting / 課題設定

Two practice videos may contain the same or highly similar performance segment but start at different times. The goal is to estimate a temporal offset (start alignment) for synchronized side-by-side playback on mobile devices, robust to noise such as:

- different camera distance/zoom and minor viewpoint changes,
- partial occlusions and low-confidence keypoints,
- timing differences (pauses, warm-up motions, idle segments),
- pose-estimation noise and missing keypoints.

---

## 2. Inputs and Outputs / 入出力

### Inputs / 入力
- Video A, Video B
- Parameters (examples; not limiting):
  - `F`: sampling FPS (frames per second)
  - `searchRangeSec`: maximum offset search range (±, seconds)
  - `τ_kp`: keypoint confidence threshold
  - `M_shared`: minimum number of shared-valid keypoints for a frame-pair to be scored (e.g., ≥2; implementation-defined)
  - `τ_sim`: similarity threshold for contiguous match detection
  - `N`: minimum contiguous frames above threshold
  - `τ_conf`: acceptance threshold for confidence
  - Hybrid parameters (optional):
    - `k`: top-K audio candidates
```

  - `minSepSec`: minimum separation between candidate peaks (seconds)
  - `r_frames` / `r_sec`: local neighborhood radius around candidates (frames / seconds)
  - `ε_frames` / `ε_sec`: micro-refinement radius for audio verification (frames / seconds)
- Optional runner-up separation (optional):
  - `δ_excl_frames` / `δ_excl_sec`: exclusion radius around Δ* when selecting runner-up

Outputs / 出力
- `startA_sec`, `startB_sec`: aligned start timestamps for synchronized playback
- `confidence`: acceptance score (e.g., normalized to [0,1])
- Optional diagnostics:
  - `bestOffset_frames`
  - `bestOffset_sec`
  - `bestSim`
  - `secondBestSim`
  - `margin`
  - `validRate`

Units / 単位（重要）

Offsets and neighborhoods can be represented in either frames or seconds. In this note, **Δ is expressed in frames by default** (integer frame shift), and its seconds representation is:

```text
Δ_sec = Δ / F
r_sec = r_frames / F
ε_sec = ε_frames / F
```

Equivalent embodiments may instead express all quantities directly in seconds.

---

3. Method / 手法

3.1 Frame Sampling / フレーム抽出

1. Decode each video into frames sampled at fixed FPS `F` and a fixed target size.
2. Frame index `t` corresponds to timestamp `t/F`.

3.2 Pose Estimation / 骨格推定（Single-pose / Multi-pose）

For each frame, estimate pose keypoints with coordinates and confidence `(x_i, y_i, s_i)`.

- **Single-pose:** directly use the single-person output.
- **Multi-pose:** select a target person from multiple detected persons using one or more of:
  - detection score / instance score,
  - average keypoint confidence,
  - spatial consistency (e.g., torso center continuity, size constraints),
  - track continuity over time (simple tracking by nearest center / IoU / embedding),
  - region-of-interest constraints,
  - rejecting implausible skeleton geometry.

This note is not limited to any specific model or framework; any on-device pose estimator producing keypoints with confidence qualifies as an equivalent embodiment.

3.3 Pose Normalization & Vectorization / 正規化ベクトル化

Purpose: reduce variation due to camera translation and scale.

- Keypoint filtering: `K_valid(t) = { i | s_i(t) ≥ τ_kp }`
- Center selection (fallback): pelvis/hip center → shoulder center → centroid of valid keypoints
- Scale selection (fallback): shoulder width → hip width → max radius from center
- Normalized coordinates (for valid keypoints):
  - `p_i'(t) = ( (x_i(t) - c_x(t)) / α(t), (y_i(t) - c_y(t)) / α(t) )`
- Feature vector: concatenate normalized coordinates in a fixed keypoint order

Missing keypoints may be handled by masking (shared-valid only), zero-imputation with mask, or variable-length aggregation normalized by count.

3.4 Offset Scan & Score / オフセット走査とスコア

Offset sign convention: For frame index `t` in A, pair with `(t + Δ)` in B. Thus, **Δ > 0 means B lags A**, and **Δ < 0 means B leads A**.

Search range (frames):
- `Δ ∈ [-searchRangeSec·F, +searchRangeSec·F]`

Discrete-range note (implementation-defined): since `searchRangeSec·F` may not be an integer, the conversion from seconds to frame bounds may apply rounding such as `round`, `floor`, or `ceil`, depending on implementation.

Shared-valid minimum:
- compute `sharedValidCount(t,Δ)` = number of keypoints valid in both frames
- exclude frame pairs with `sharedValidCount(t,Δ) < M_shared`
- `M_shared` is implementation-defined (e.g., ≥2) and tunable

Per-frame similarity:
- example: `sim(t,Δ) = cos( vA(t), vB(t+Δ) )`
- distance form: `dist(t,Δ) = 1 - sim(t,Δ)`

Offset score:
- `Score(Δ) = average_t sim(t,Δ)` over included pairs (or minimize average distance)

Best and runner-up:
- `Δ* = argmax Score(Δ)`
- `Δ2 = argmax_{Δ≠Δ*} Score(Δ)`
- optional runner-up separation: exclude `|Δ - Δ*| ≤ δ_excl_frames` to avoid near-duplicate peaks

Optional (non-limiting): enforce a minimum overlap length and discard too-short overlaps.

3.5 Robust Start Point via Contiguous Match / 連続一致区間による開始点
- Find earliest `t0` such that `sim(t,Δ*) ≥ τ_sim` holds for at least `N` consecutive frames (valid frame pairs only).
- If not found, staged relaxation retries may lower `τ_sim`, reduce `N`, switch to window-ratio criteria, switch to distance thresholds, etc.

Aligned start timestamps:
- Global search:
  - `startA_sec = t0 / F`
  - `startB_sec = (t0 + Δ*) / F`
- Windowed/local search (seeded neighborhoods, Phase1 restricted windows):
  - `startA_sec = baseA_sec + t0 / F`
  - `startB_sec = baseB_sec + (t0 + Δ*) / F`

3.6 Confidence & Acceptance / 信頼度と採否
Definitions:
- `bestSim = Score(Δ*)`
- `secondBestSim = Score(Δ2)`
- `margin = bestSim - secondBestSim`

Validity ratio `validRate` (over matched segment/window):
- `validFrame(t,Δ)` = 1 if the pair is eligible and included (e.g., `sharedValidCount≥M_shared`), else 0
- `validFrameRate` = average of `validFrame(t,Δ*)`
- optional: include validKeypointRate; validRate may denote valid-frame ratio and/or a combination

Confidence (example):
```text
confidence = 0.5 × bestSim + 0.3 × margin + 0.2 × validRate
```

Accept if `confidence ≥ τ_conf`, else reject.

Distance-minimization note:

- if computed via distance minimization internally, `bestSim`/`secondBestSim` for margin may be computed after any monotonic transform (e.g., distance→similarity)
- values may be clipped/normalized to [0,1] prior to thresholding

---

4. Optional Hybrids (Audio Assistance and/or Verification) / オプション：音声支援の粗候補生成および／または音声検証・微調整

4.1 Audio-based coarse candidate generation (optional) / 音声で粗候補生成（任意）
Compute audio similarity using any audio features (energy envelope, onset patterns, spectral features, fingerprint, cross-correlation, etc.). Generate top-K candidate offsets (in frames or seconds).

Optional robustness:
- enforce minimum separation between candidate peaks `minSepSec` to avoid near-duplicate candidates,
- keep both positive and negative candidates (respecting Δ sign convention).

4.2 Pose refinement around candidates (optional) / Poseで候補近傍を精密化（任意）
For each candidate, refine only within a local neighborhood:
- `Δ ∈ [Δc_i - r_frames, Δc_i + r_frames]`
  (or seconds: `Δ_sec ∈ [Δc_sec - r_sec, Δc_sec + r_sec]`)

Compute `Score(Δ)` using pose matching (Section 3), select best `Δ*` across all neighborhoods, then apply contiguous-match and confidence.

4.3 Post-pose audio verification / micro refinement (optional) / Pose後の音声検証・微調整（任意）
After a pose-derived offset `Δ*` is selected:
- evaluate audio similarity in a small neighborhood `Δ ∈ [Δ* - ε_frames, Δ* + ε_frames]` (or seconds)
- if a nearby offset yields better audio consistency (or passes an audio check), adjust

`Δ*` accordingly
- or reject if audio strongly contradicts the pose-derived alignment

---

5. Optional Pose-only Phase Search Acceleration (Phase1; example) / Poseのみの探索高速化（Phase1；例示）

To improve runtime on mobile devices, an optional pose-only acceleration phase can be used:
- **Early-segment restricted search:** search only within early segments (e.g., 0–10s, 10–20s) before expanding.
- **High-motion window selection:** detect windows with high motion magnitude and search only those windows. Examples of motion proxies include pose displacement norms across frames, simple frame-difference energy, optical-flow statistics (if available), and other lightweight motion proxies.
- **Fallback:** if no confident match is found in restricted windows, fall back to the full-range scan (Section 3.4).

This phase acceleration is orthogonal to audio-assisted hybrids and can be combined with them.

---

6. Failure Handling and Fallbacks / 失敗時のフォールバック
If pose alignment fails or `confidence < τ_conf`:
- fall back to audio-only alignment (if available), and/or
- keep the pose offset but require audio verification before acceptance, and/or
- reject automatic alignment and allow manual adjustment.

---

7. Publication Metadata / 公開メタ情報
For a fixed release, record DOI and immutable URLs in external metadata (README / GitHub Release / Zenodo). If you also include them here, ensure you treat this file as

versioned content.

- Authors / 著者: Sparklyforce
- Version: v1.0.1
- First public timestamp (UTC, ISO 8601): YYYY-MM-DDThh:mm:ssZ
- Repository tag: v1.0.1
- Commit (Git SHA): <to be filled>
- GitHub Release (Permanent URL): <to be filled>
- Zenodo record (Immutable archive URL): <to be filled>
- Zenodo Version DOI: <to be filled>
- Zenodo Concept DOI (all versions): <to be filled>

License / ライセンス